

```
# libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression,Lasso
from sklearn.metrics import accuracy_score,mean_squared_error, mean_absolute_error, r2_score
```

```
housing = pd.read_csv('Housing.csv')
```

```
housing
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking
0	13300000	7420	4	2	3	yes	no	no	no	yes	2
1	12250000	8960	4	4	4	yes	no	no	no	yes	3
2	12250000	9960	3	2	2	yes	no	yes	no	no	2
3	12215000	7500	4	2	2	yes	no	yes	no	yes	3
4	11410000	7420	4	1	2	yes	yes	yes	no	yes	2
...	...	...	...	...	...	...	...	...	...	...	...
540	1820000	3000	2	1	1	yes	no	yes	no	no	2
541	1767150	2400	3	1	1	no	no	no	no	no	0
542	1750000	3620	2	1	1	yes	no	no	no	no	0
543	1750000	2910	3	1	1	no	no	no	no	no	0
544	1750000	3850	3	1	2	yes	no	no	no	no	0

545 rows × 13 columns

```
housing.isnull().sum()
```

	0
price	0
area	0
bedrooms	0
bathrooms	0
stories	0
mainroad	0
guestroom	0
basement	0
hotwaterheating	0
airconditioning	0
parking	0
prefarea	0
furnishingstatus	0

dtype: int64

```
housing.duplicated().sum()
```

```
np.int64(0)
```

```
housing.columns = housing.columns.str.strip()
```

```
# Set up the figure
fig, axes = plt.subplots(nrows=5, ncols=3, figsize=(30, 20))
axes = axes.flatten()

sns.histplot(housing['price'], ax=axes[0], color='green')
```

```

sns.histplot(housing['area'], ax=axes[1],color='red')

sns.histplot(housing['bedrooms'], ax=axes[2],color='yellow')

sns.histplot(housing['bathrooms'], ax=axes[3],color='blue')

sns.histplot(housing['stories'], ax=axes[4],color='red')

sns.countplot(x=housing['mainroad'], ax=axes[5],color='yellow')

sns.countplot(x=housing['guestroom'], ax=axes[6],color='blue')

sns.countplot(x=housing['basement'], ax=axes[7],color='red')

sns.countplot(x=housing['hotwaterheating'], ax=axes[8],color='yellow')

sns.countplot(x=housing['airconditioning'], ax=axes[9],color='blue')

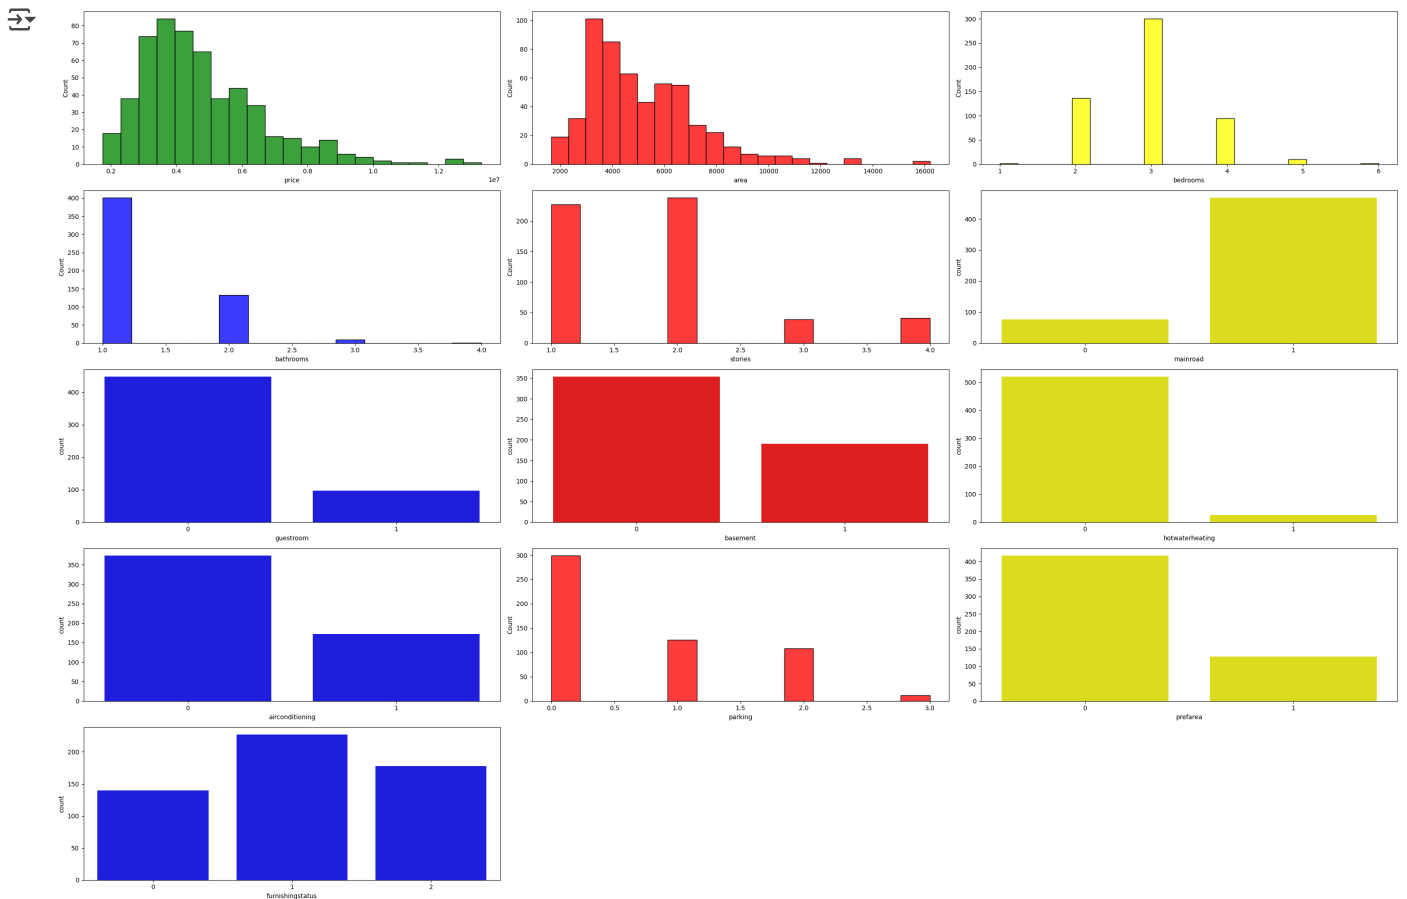
sns.histplot(housing['parking'], ax=axes[10],color="red")

sns.countplot(x=housing['prefarea'], ax=axes[11],color="yellow")

sns.countplot(x=housing['furnishingstatus'], ax=axes[12],color="blue")

for j in range(13, len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()


```




```

plt.figure(figsize=(30, 10))
sns.lineplot(x=housing['area'], y=housing['price'], color='blue')
plt.title('Price vs Area')
plt.xlabel('Area')
plt.ylabel('Price')
plt.grid(True)
plt.show()

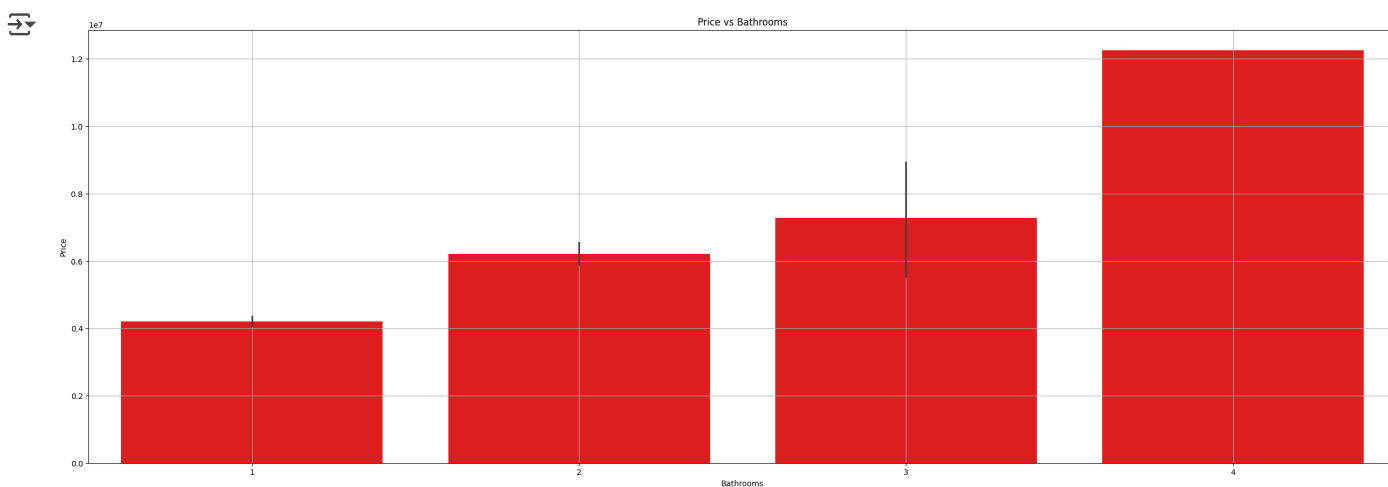
```



```
plt.figure(figsize=(30, 5))
plt.scatter(housing['bedrooms'], housing['price'], color='green')
plt.title('Price vs Bedrooms')
plt.xlabel('Bedrooms')
plt.ylabel('Price')
plt.grid(True)
plt.show()
```

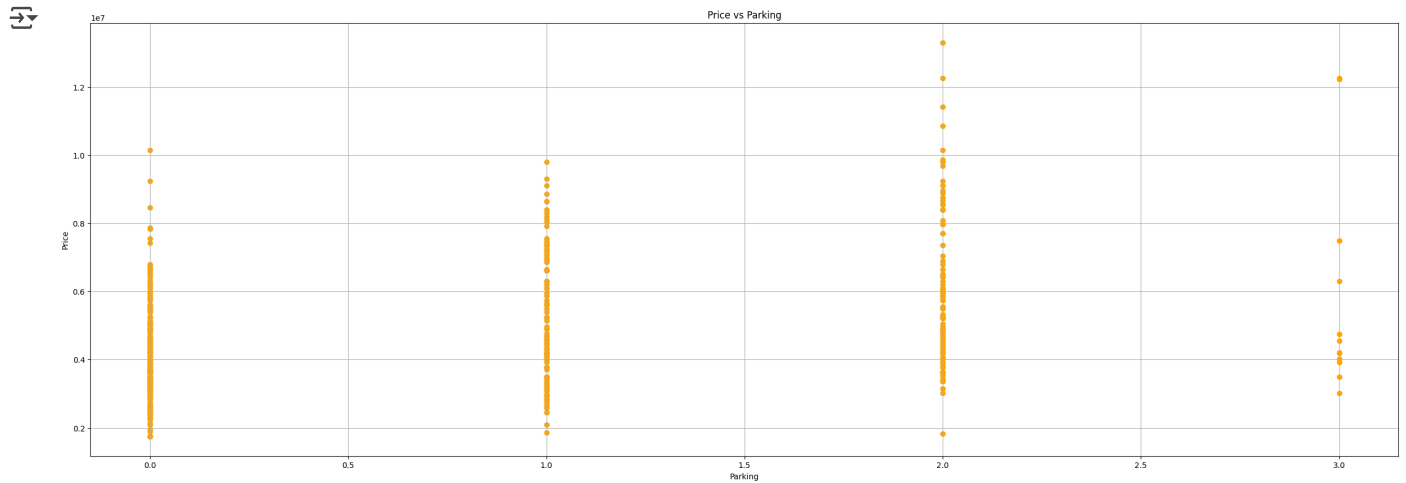


```
plt.figure(figsize=(30, 10))
sns.barplot(x=housing['bathrooms'], y=housing['price'], color='red')
plt.title('Price vs Bathrooms')
plt.xlabel('Bathrooms')
plt.ylabel('Price')
plt.grid(True)
plt.show()
```



```
plt.figure(figsize=(30, 10))
plt.scatter(housing['parking'], housing['price'], color='orange')
plt.title('Price vs Parking')
plt.xlabel('Parking')
plt.ylabel('Price')
```

```
plt.grid(True)
plt.show()
```



```
#Encoding categorical features
le = LabelEncoder()
housing['mainroad'] = le.fit_transform(housing['mainroad'])
housing['guestroom'] = le.fit_transform(housing['guestroom'])
housing['basement'] = le.fit_transform(housing['basement'])
housing['hotwaterheating'] = le.fit_transform(housing['hotwaterheating'])
housing['airconditioning'] = le.fit_transform(housing['airconditioning'])
housing['prefarea'] = le.fit_transform(housing['prefarea'])
housing['furnishingstatus'] = le.fit_transform(housing['furnishingstatus'])
#creating new features
housing['total_rooms'] = housing['bedrooms'] + housing['bathrooms']
housing['price_per_room'] = housing['price'] / housing['total_rooms']
housing['price_per_area'] = housing['price'] / housing['area']
housing['price_per_bathroom'] = housing['price'] / housing['bathrooms']
housing['area_bedrooms'] = housing['area'] * housing['bedrooms']
housing['area_bathrooms'] = housing['area'] * housing['bathrooms']
housing['area_furnishingstatus'] = housing['area'] * housing['furnishingstatus']

# deal with skewed price and area
housing['log_price'] = np.log(housing['price'])
housing['log_area'] = np.log(housing['area'])
```

```
housing
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	...	fur
0	13300000	7420	4	2	3	1	0	0	0	1	...	
1	12250000	8960	4	4	4	1	0	0	0	1	...	
2	12250000	9960	3	2	2	1	0	1	0	0	...	
3	12215000	7500	4	2	2	1	0	1	0	1	...	
4	11410000	7420	4	1	2	1	1	1	0	1	...	
...	...	...	...	...	...	...	...	...	...	...	...	
540	1820000	3000	2	1	1	1	0	1	0	0	...	
541	1767150	2400	3	1	1	0	0	0	0	0	...	
542	1750000	3620	2	1	1	1	0	0	0	0	...	
543	1750000	2910	3	1	1	0	0	0	0	0	...	
544	1750000	3850	3	1	2	1	0	0	0	0	...	


545 rows x 22 columns

```
correlation = housing.corr()['price'].sort_values(ascending=False)
print(correlation)
fig, ax = plt.subplots(figsize=(20, 15))
sns.heatmap(housing.corr(), annot=True, cmap='coolwarm')
plt.show()
```

```

price 1.000000
log_price 0.969010
price_per_room 0.768193
area_bathrooms 0.674764
area_bedrooms 0.611196
price_per_bathroom 0.570693
log_area 0.561231
area 0.535997
bathrooms 0.517545
total_rooms 0.511809
airconditioning 0.452954
stories 0.420712
price_per_area 0.392921
parking 0.384394
bedrooms 0.366494
prefarea 0.329777
mainroad 0.296898
guestroom 0.255517
basement 0.187057
hotwaterheating 0.093073
area_furnishingstatus -0.013093
furnishingstatus -0.304721
Name: price, dtype: float64

```



```

#split features and target
X = housing.drop(['price','log_price'],axis=1)
y = housing['log_price']

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

model = LinearRegression()
model.fit(X_train, y_train)

```

LinearRegression ⓘ ?  
LinearRegression()

```
y_pred = model.predict(X_test)

accuracy_score=model.score(X_test,y_test)
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
print('----- Mohamed-----')
print(' ')
print(accuracy_score*100)
print(' ')
print('-----Ayman-----')
print(' ')
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(' ')
print('-----Fathy-----')
print(' ')
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
print(' ')
print('-----EL-Halmoushy-----')
print(' ')
print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(' ')
print('-----')

print(f"R2 Score: {r2:.4f}")
```

```
----- Mohamed-----
98.11790745806447

-----Ayman-----

Mean Squared Error (MSE): 0.0036

-----Fathy-----

Root Mean Squared Error (RMSE): 0.0603

-----EL-Halmoushy-----


Mean Absolute Error (MAE): 0.0429

-----
R2 Score: 0.9812
```

```
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(25, 10))

axes[0].scatter(y_test, y_pred, color='red')
axes[0].set_xlabel("Actual Log Price")
axes[0].set_ylabel("Predicted Log Price")
axes[0].set_title("Actual vs Predicted (Scatter Plot)")
axes[0].grid(True)
sns.lineplot(x=y_test, y=y_pred, color='blue', ax=axes[0])
axes[0].set_xlabel("Actual Log Price")
axes[0].set_ylabel("Predicted Log Price")
axes[0].set_title("Actual vs Predicted (Line Plot)")
axes[0].grid(True)
for j in range(1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



```
lasso = Lasso(alpha=0.01)
lasso.fit(X_train, y_train)
y_pred_lasso = lasso.predict(X_test)
mse = mean_squared_error(y_test, y_pred_lasso)
mae = mean_absolute_error(y_test, y_pred_lasso)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred_lasso)
acc_lasso=lasso.score(X_test,y_test)
print(acc_lasso*100)
print(f"Lasso Regression Results:")
print(f"R² Score: {r2:.4f}")
print(f"RMSE: {rmse:.4f}")
print(f"MAE: {mae:.4f}")
```

```
95.75936142598209
Lasso Regression Results:
R² Score: 0.9576
RMSE: 0.0905
MAE: 0.0636
```

Start coding or [generate](#) with AI.